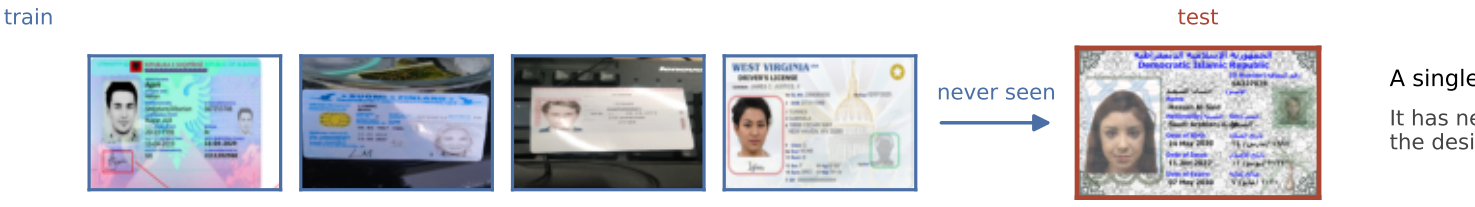
A. The gap: detectors score near-perfectly on the designs they were trained on, and fall on new ones



within the training designs: 96 to 100 % accurate

on designs not in training: 51 to 71 % (IDNet, West Virginia to Arizona); 82 to 91 % (SIDTD, five-shot); to miss only 5 % of attacks, a published detector refers 62 % of genuine cards (FantasyID)

B. This study tests on designs, a source, a capture and a script the detector has never seen



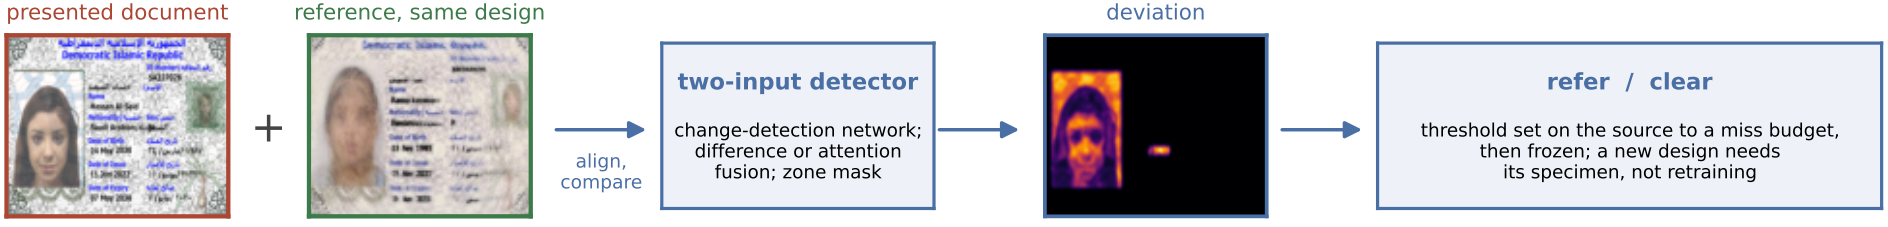
SIDTD (10 designs) and IDNet (Arizona, West Virginia): text and face alterations

FantasyID, never trained on: printed and photographed (an Arabic card altered by us for illustration)

A single-image detector sees only this image.

It has never seen the design, so it cannot tell the design's own features from an alteration.

C. The method: give the detector the specimen of the same design, and score the deviation



from the reader

design median, looked up by the design the reader identifies

where the presented page departs from its design

D. What is reported, in the first line's own terms

with and without the specimen

same network, same data: how much of the loss does it recover?

referral at a miss budget

threshold carried to unseen designs: realized misses and referrals

three protocols

unseen design; unseen source and capture; the published IDNet pair

time on a CPU

per document, batch size one, with cached reference features